

Passwords, Logins, and CHF



Ashkan Jangodaz
British Columbia Institute of Technology

The Real Problem

- Most security breaches don't involve sophisticated hacking.
- Over 80% of breaches involve stolen or weak credentials. [1]
 - If someone has your password, they can simply *be* you.
- Authentication answers one question:
 - *Who are you?*
- The problem:
 - Secrets can be copied, shared, stolen, or guessed.

Factors of Authentication

| Factor | Description | Examples |
|--------------------|-----------------------------|-----------------------|
| Something you know | Secret in your memory | Password, PIN |
| Something you have | Physical object you possess | Phone, hardware token |
| Something you are | Biometric characteristic | Fingerprint, face |

How Should Systems Store Passwords?

- **Wrong:** Store passwords in plain text
 - Database stolen → all passwords immediately exposed
- **Wrong:** Encrypt passwords
 - Encryption is reversible
 - If attacker gets the key, all passwords exposed
- **Correct:** Store password *hashes*
 - One-way transformation
 - Cannot be reversed
 - System never needs to store actual password

Password Verification

- **When you set a password:**
 - **You type:** MySecretPassword
 - **System computes:** hash("MySecretPassword") → a7f3b2c9...
 - **System stores:** a7f3b2c9...
 - Original password discarded
- **When you log in:**
 - You type: MySecretPassword
 - System computes hash of what you typed
 - Compares to stored hash
 - Match → access granted

What Linux Actually Stores

| File | Purpose | Readable By |
|-------------|---|--------------------|
| /etc/passwd | Identity information (username, UID, home, shell) | Everyone |
| /etc/shadow | Password verification data (hashes) | Root only |

Anatomy of /etc/shadow

- **alice:\$y\$j9T\$7R3xK9pL\$WvM8nQ2kF5hJ1xY4bN6cR9dG3mP7sTO
wA2zX8vB4:19750:0:99999:7:30:20089:**

| Field | Value | Meaning |
|---------------|---------------------------------|--------------------------------------|
| Username | alice | Account name |
| Password hash | \$y\$j9T\$7R3xK9pL\$WvM8nQ2k... | yescrypt hash with salt |
| Last changed | 19750 | Days since Jan 1, 1970 |
| Min age | 0 | Can change password anytime |
| Max age | 99999 | Password never expires (~274 years) |
| Warn period | 7 | Warn user 7 days before expiry |
| Inactive | 30 | Disable account 30 days after expiry |
| Expire date | 20089 | Account expires on Jan 1, 2025 |

The Hash Format

- **\$id\$salt\$hash**

| Component | Purpose |
|-----------|--------------------------------------|
| \$id\$ | Identifies the hashing algorithm |
| salt | Random value unique to this password |
| hash | The actual hash output |

- **Common identifiers:**

- \$1\$ → MD5crypt (legacy, weak)
- \$5\$ → SHA-256crypt
- \$6\$ → SHA-512crypt
- \$y\$ → yescrypt (modern Linux default)
- \$2a\$, \$2y\$ → bcrypt

What is a Salt?

- **Problem without salt:**

- Two users with password "password123" get identical hashes
- Attacker can precompute hashes for common passwords (rainbow tables)
- Attacker can see which users have the same password

- **With salt:**

- Each password gets a unique random salt
- Same password + different salts → different hashes
- Precomputation becomes useless

Important About Salts

- Salts are stored alongside the hash
 - Visible in /etc/shadow
 - This is intentional, system needs salt to verify passwords
- Salts prevent pre-computation attacks (rainbow tables)
- Salts do NOT prevent brute-force attacks
- Attacker with hash file has all the salts too
- **Salts ensure each password must be cracked independently.**

Dictionary Attacks

- Attackers don't start with "aaaaaa" and work up, they start with lists of known passwords:
 - Passwords from previous breaches
 - Common words and phrases
 - Keyboard patterns (qwerty, 123456)
 - Common substitutions (p@ssw0rd, passw0rd)
- If your password is in a wordlist, length doesn't matter.

Lab

- Separation between /etc/passwd and /etc/shadow
- Stored value is verification material, not the password
- If dictionary cracking succeeds quickly, weakness is predictability
- Dictionary attacks exploit human choice
- Brute force exploits small search space
- Strong rules don't help if password is predictable

Lab

- Create a test account with a password you choose
- Examine how Linux stores identity vs. password data
- Extract a hash and understand its structure
- Attempt dictionary cracking on your own password
- Compare dictionary attack to brute force
- Clean up completely when finished

Hash Functions

► A **hash function H** accepts a variable-length **block of data M** as input and produces a **fixed-size hash value (hash code) h**

- $h = H(M)$
- Principal object is data integrity

► **Cryptographic hash function**

- An algorithm for which it is computationally infeasible to find either:
 - a data object that maps to a pre-specified hash result (the one-way property)
 - two data objects that map to the same hash result (the collision-free property)



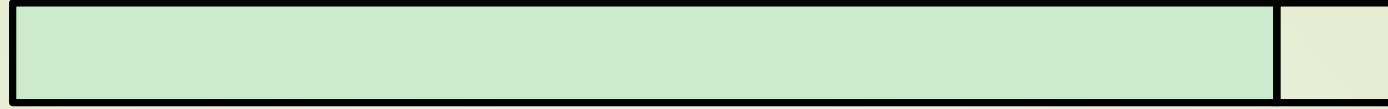
Hash Functions

- ▶ A change to any bit or bits in **M** results, with high probability, in a change to the hash value **h**
 - Hash functions are often used to determine whether or not data has changed
- ▶ Cryptographic hash functions are **keyless** cryptographic algorithms
 - There are **however** some “keyed” hash functions (e.g. MAC)



Hash Functions

- ▶ Typically, the input is **padded** out to an integer multiple of some fixed length (e.g., 1024 bits)
 - Padding includes the value of the length of the original message in bits.
 - The **length field** is a security measure to increase the difficulty for an attacker to produce an alternative message with the same hash value





Hash Functions Applications

► Message Authentication

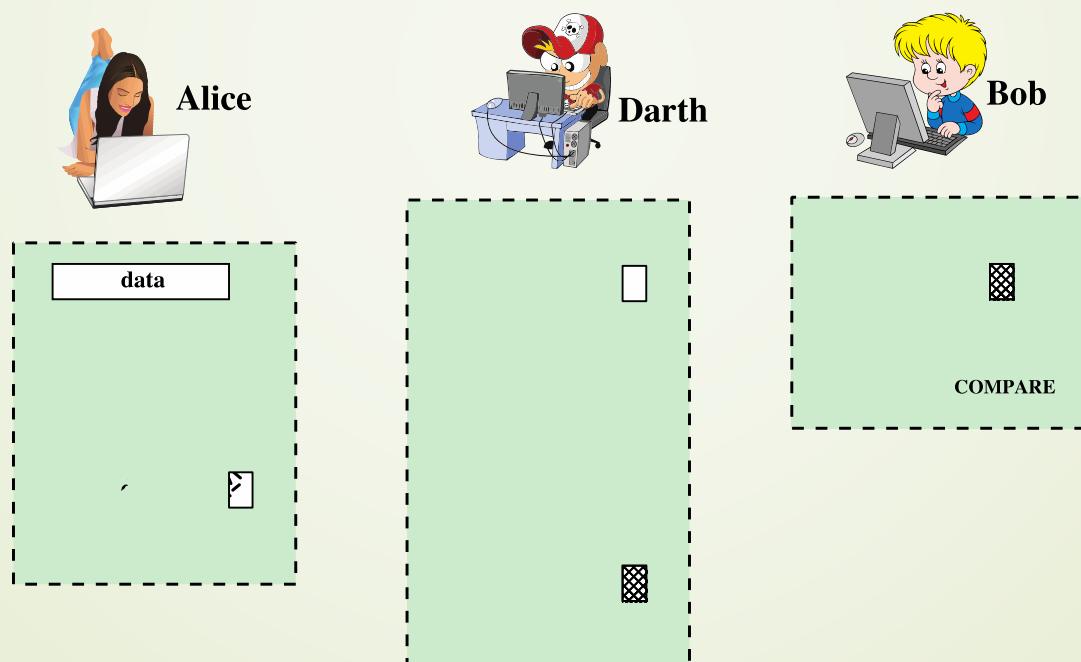
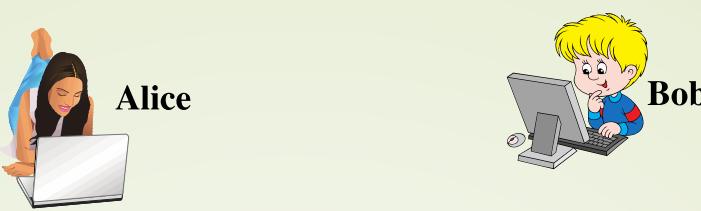
- A mechanism or service used to verify the **integrity** of a message
- Message authentication assures that data received are exactly as sent (i.e., there is no modification, insertion, deletion, or replay)
- The hash value is often referred to as a **message digest**
- Is achieved using a Message Authentication code (MAC)

Message Authentication Code (MAC)

- ▶ Also known as a **keyed** hash function
- ▶ Typically used between two parties that share a secret key to authenticate information exchanged between those parties

Takes as input a secret key and a data block and produces a hash value (MAC) which is associated with the protected message

- If the integrity of the message needs to be checked, the MAC function can be applied to the message and the result compared with the associated MAC value
- An attacker who alters the message will be unable to alter the associated MAC value without knowledge of the secret key



(b) Man-in-the-middle attack

Figure 11.2 Attack Against Hash Function

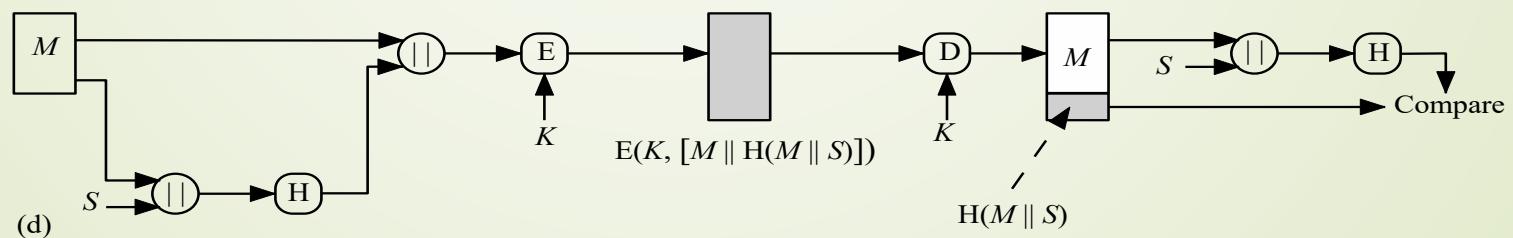
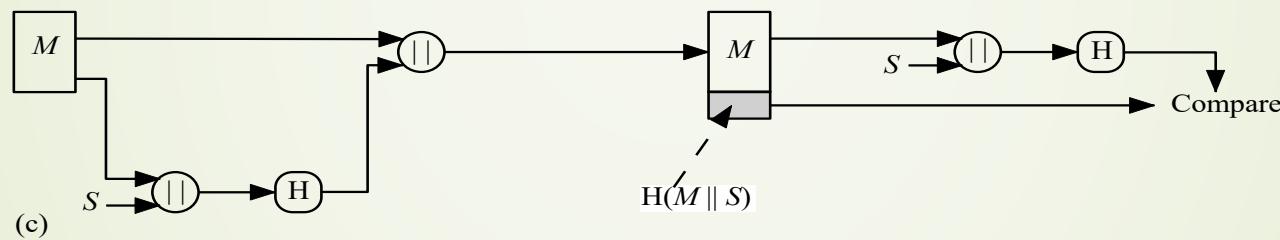
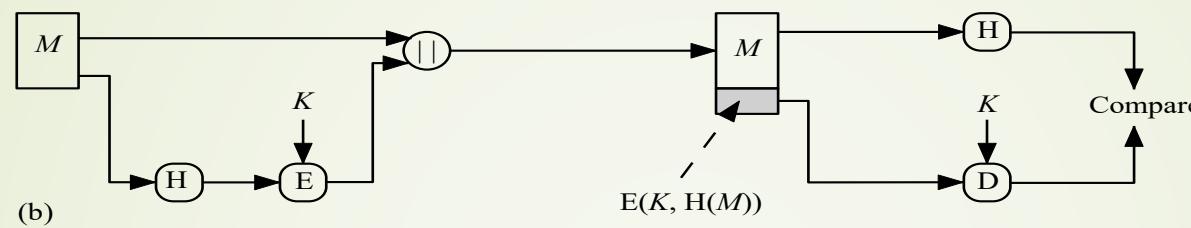
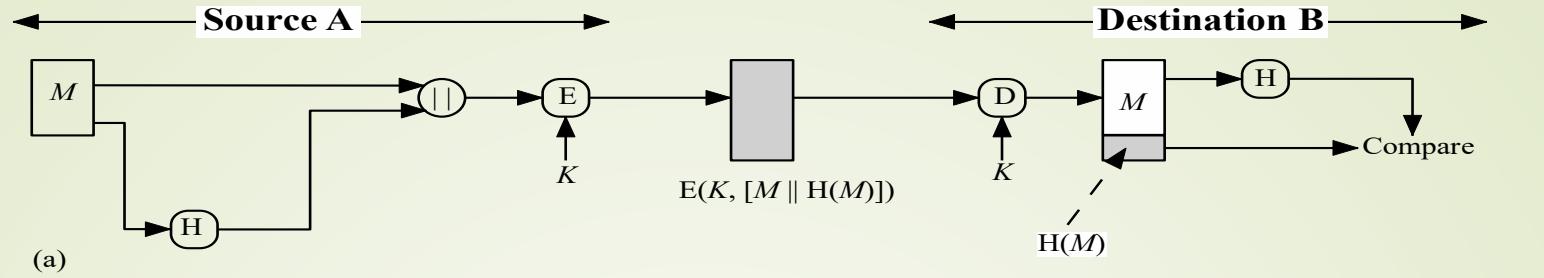


Figure 11.3 Simplified Example of the use of a Hash Function for Message Authentication



Hash Functions Applications

► Digital Signature

- Operation is similar to that of the MAC
- The hash value of a message is encrypted with a user's private key
- Anyone who knows the user's public key can verify the integrity of the message
- An attacker who wishes to alter the message would need to know the user's private key
- Implications of digital signatures go beyond just message authentication

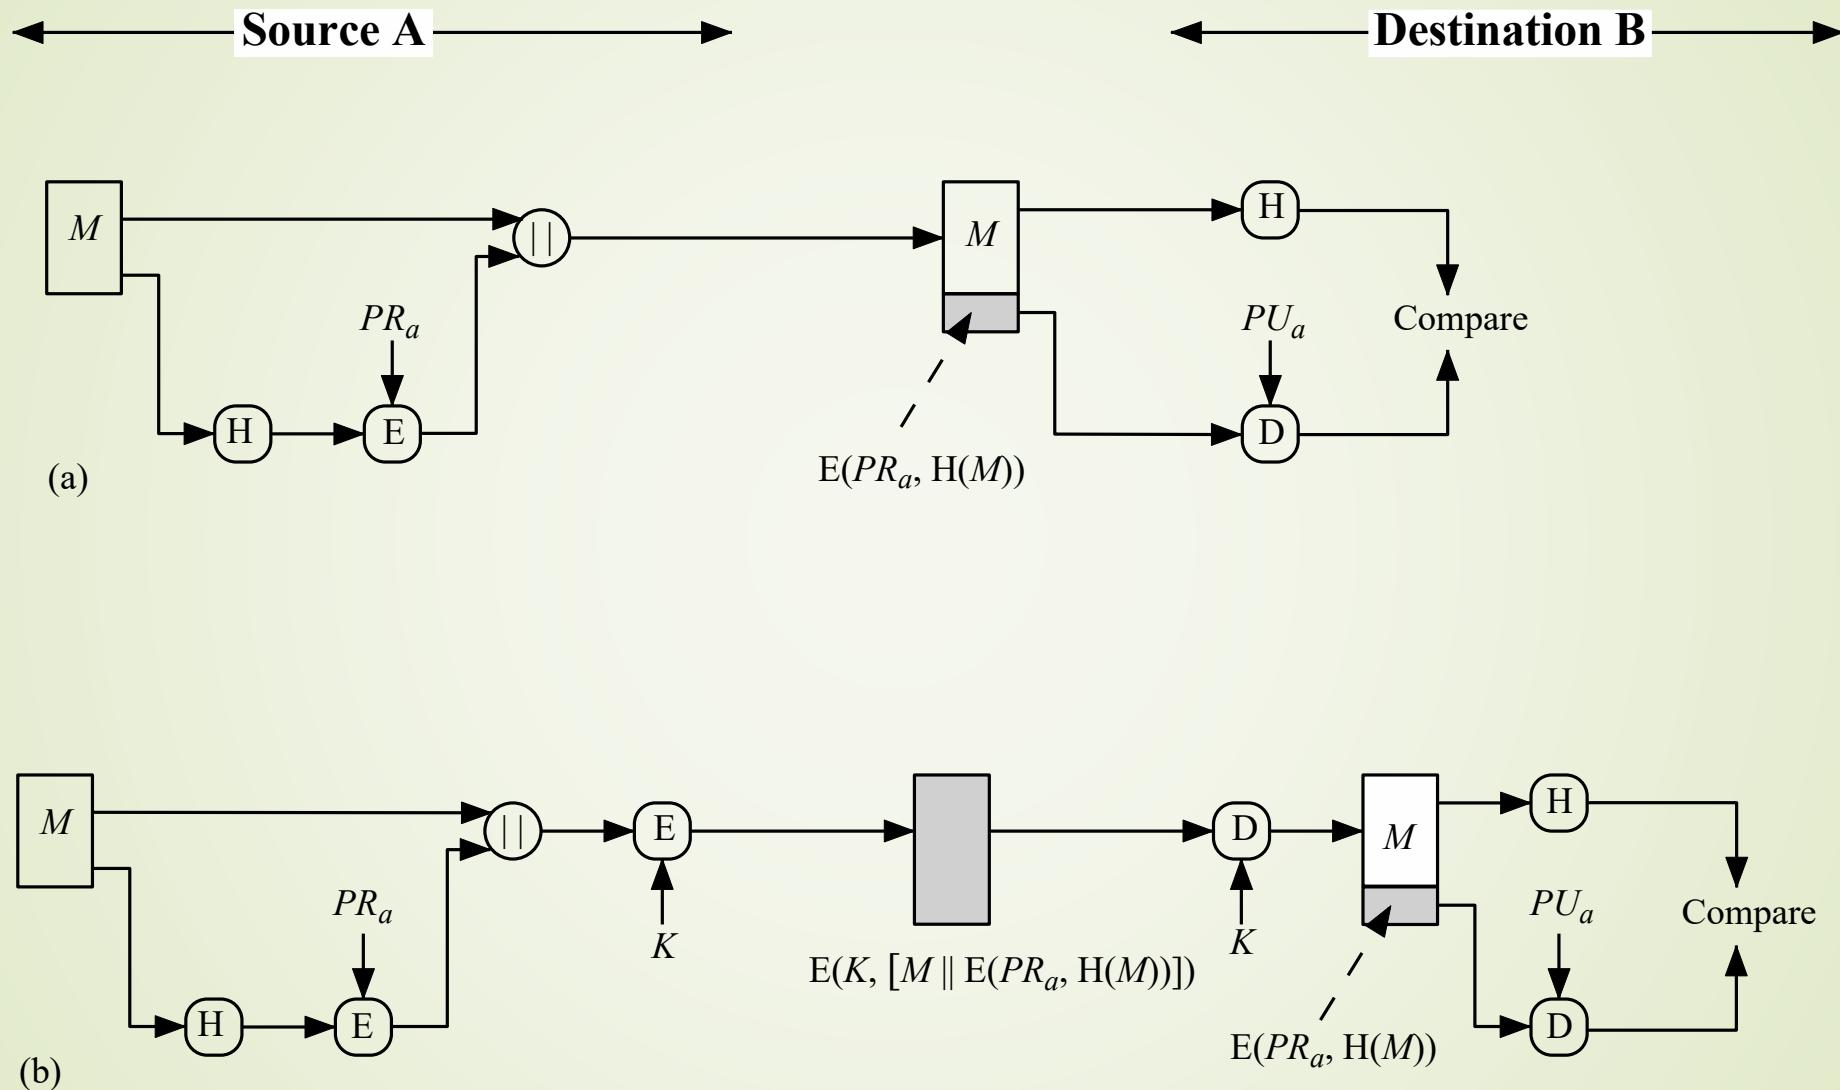
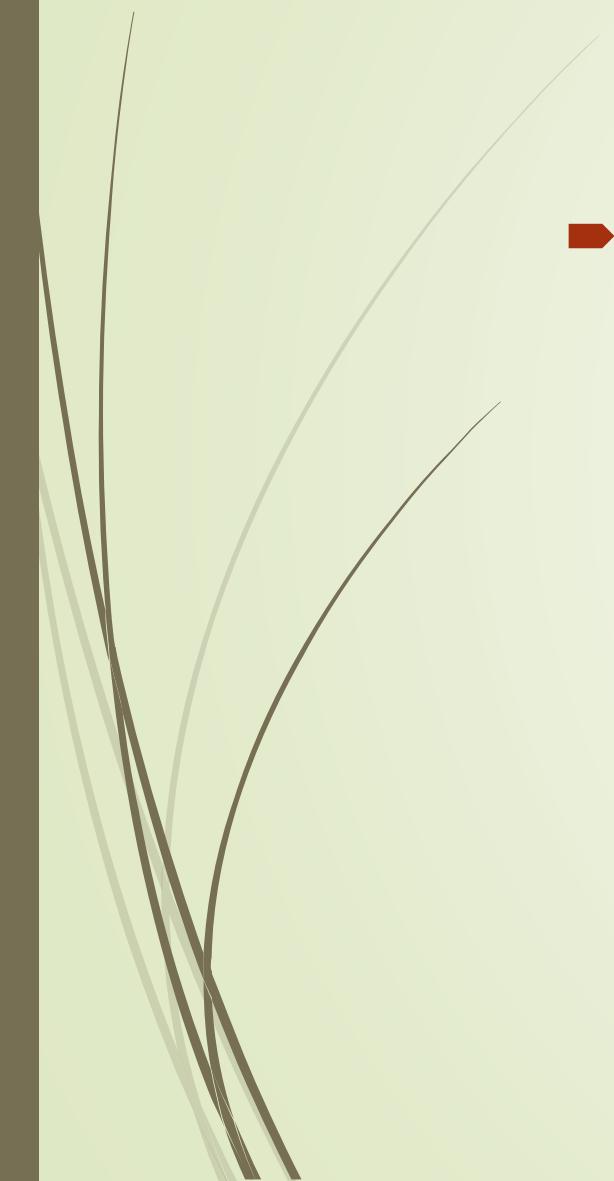


Figure 11.4 Simplified Examples of Digital Signatures



Other Hash Functions Applications

- ▶ Commonly used to **create a one-way password file**
 - When a user enters a password, the hash of that password is compared to the stored hash value for verification
 - This approach to password protection is used by most operating systems
- ▶ Can be used for **intrusion** and **virus detection**
 - Store $H(F)$ for each file on a system and secure the hash values
 - One can later determine if a file has been modified by recomputing $H(F)$
 - An intruder would need to change F without changing $H(F)$



Other Hash Functions Applications

- ▶ Can be used to **construct a pseudorandom function (PRF) or a pseudorandom number generator (PRNG)**
 - A common application for a hash-based PRF is for the generation of symmetric keys



Two Simple Hash Functions

- ▶ Consider two simple insecure hash functions that operate using the following general principles:
 - The input (message, file, etc.) is viewed as a sequence of n-bit blocks
 - The input is processed one block at a time in an iterative fashion to produce an n-bit hash function
- 1. Bit-by-bit exclusive-OR (XOR) of every block
 - $C_i = b_{i1} \oplus b_{i2} \oplus \dots \oplus b_{im}$
 - Produces a simple parity for each bit position and is known as a **longitudinal redundancy check** (p. 13 to 15)
 - Reasonably **effective for random data** as a data integrity check

Two Simple Hash Functions

- The probability that a data error will result in an unchanged hash value is 2^{-n}
 - With more predictably formatted data, the function is less effective
 - For example, in most normal text files, the high-order bit of each octet is always zero
 - if a **128-bit** hash value is used, instead of an effectiveness of 2^{-128} , the hash function on this type of data has an effectiveness of 2^{-112}
 - ASCII for 'H': **01001000**, for 'i': **01101001**, and for '!': **00100001**



Two Simple Hash Functions

2. A simple way to improve matters is to perform a one-bit circular shift, or rotation, on the hash value after each block is processed.
 - Initially set the n-bit hash value to zero.
 - Process each successive n-bit block of data as follows:
 1. Rotate the current hash value to the left by one bit.
 2. XOR the block into the hash value.
- This has the effect of “randomizing” the input more completely and overcoming any regularities that appear in the input

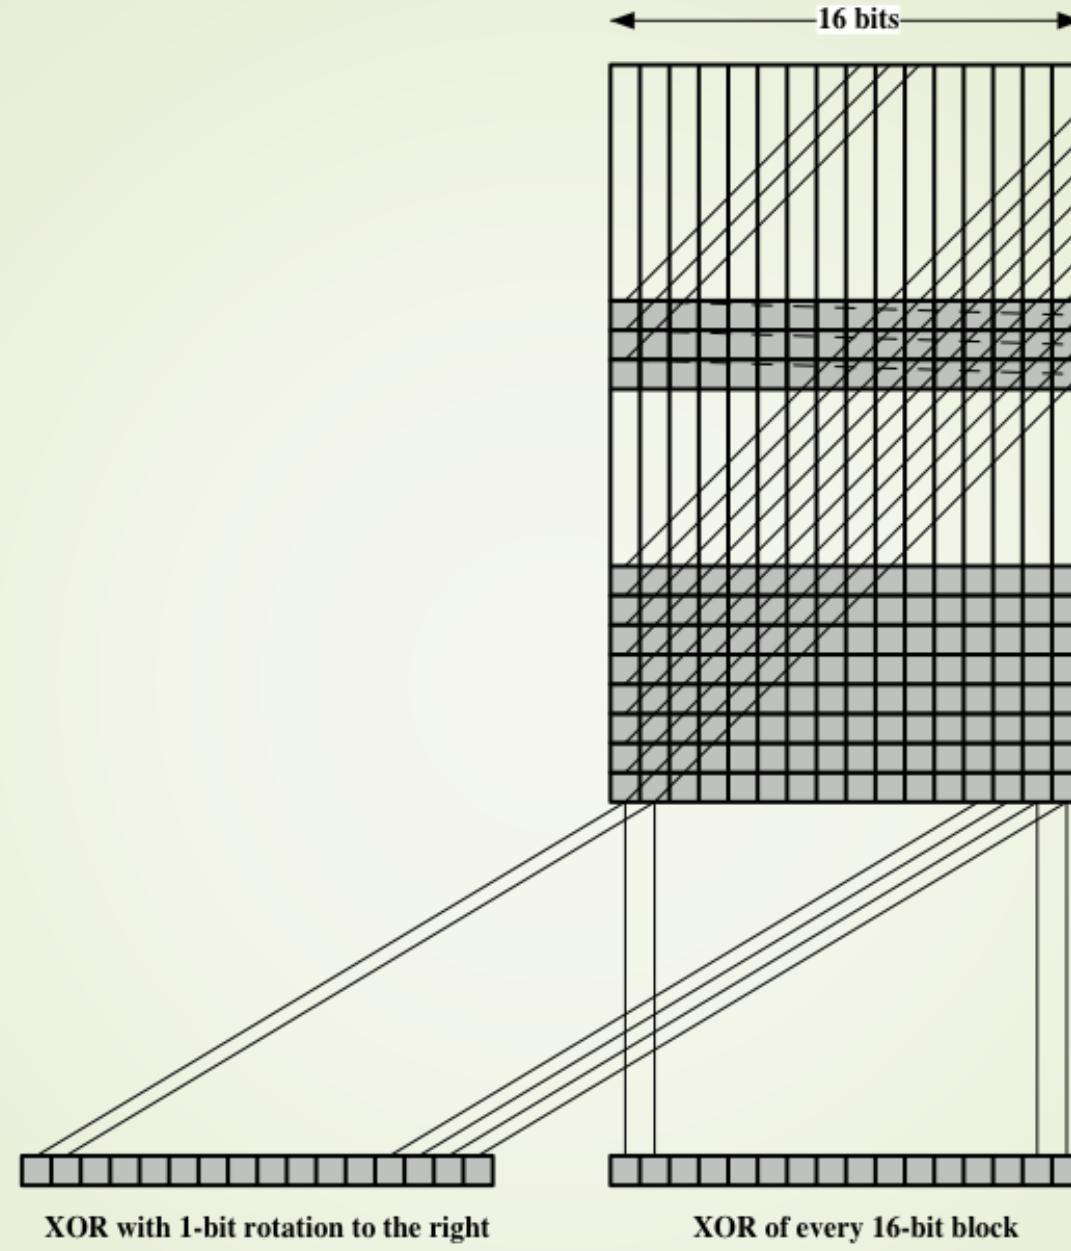


Figure 11.5 Two Simple Hash Functions (16-bit hash values)

Two Simple Hash Functions

- ▶ Although the second procedure provides a good measure of data integrity, it is virtually useless for data security when an encrypted hash code is used with a plaintext message
 - **Figures 11.3 (b) and 11.4 (a)**
- ▶ Given a message, it is an easy matter to produce a new message that yields that hash code
- ▶ A simple XOR or rotated XOR (RXOR) is insufficient if only the hash code is encrypted
 - Even if the entire message is encrypted, they are not totally sufficient (p. 345 and 346 of the textbook)

Requirements and Security

► Preimage

- x is the preimage of h for a hash value $h = H(x)$
- Is a data block whose hash function, using the function H , is h
- Because H is a many-to-one mapping, for any given hash value h , there will in general be **multiple preimages**

► Collision

- Occurs if we have $x \neq y$ and $H(x) = H(y)$
- Because we are using hash functions for data integrity, collisions are clearly undesirable



Number of Potential Collisions

- ▶ Suppose the length of the hash code is n bits
- ▶ The function H takes as input messages or data blocks of length b bits
 - $b > n$
 - The total number of possible messages is 2^b and the total number of possible hash values is 2^n
 - On average, each hash value corresponds to 2^{b-n} preimages
 - If H tends to uniformly distribute hash values, each hash value will have close to 2^{b-n} preimages
- ▶ Now imagine b to be variable size!

| Requirement | Description |
|--|---|
| Variable input size | H can be applied to a block of data of any size. |
| Fixed output size | H produces a fixed-length output. |
| Efficiency | $H(x)$ is relatively easy to compute for any given x , making both hardware and software implementations practical. |
| Preimage resistant (one-way property) | For any given hash value h , it is computationally infeasible to find y such that $H(y) = h$. |
| Second preimage resistant (weak collision resistant) | For any given block x , it is computationally infeasible to find $y \neq x$ with $H(y) = H(x)$. |
| Collision resistant (strong collision resistant) | It is computationally infeasible to find any pair (x, y) such that $H(x) = H(y)$. |
| Pseudorandomness | Output of H meets standard tests for pseudorandomness |

Table 11.1 Requirements for a Cryptographic Hash Function H

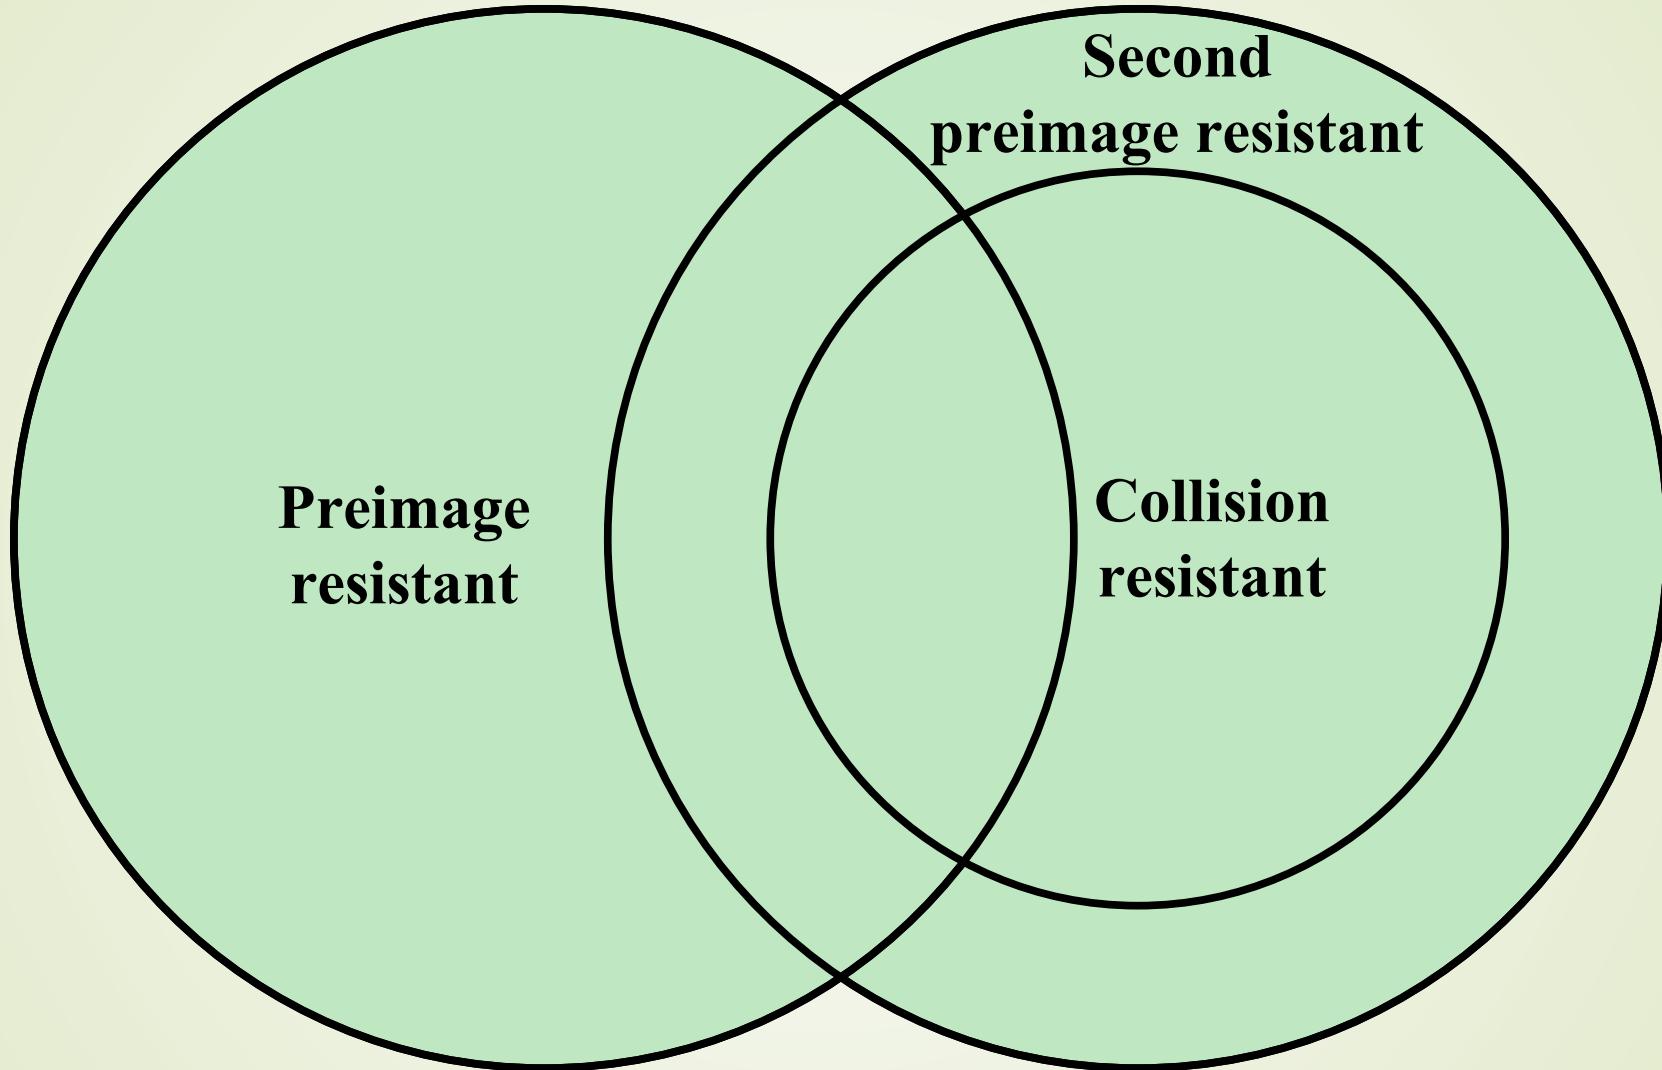
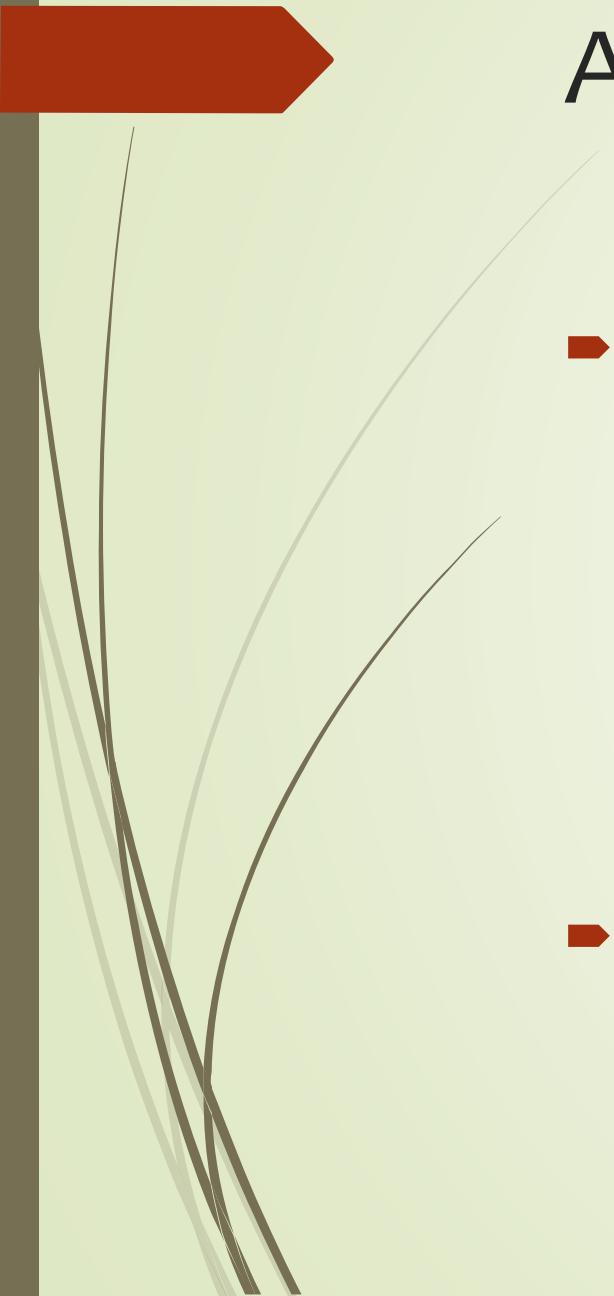


Figure 11.6 Relationship Among Hash Function Properties

| | Preimage Resistant | Second Preimage Resistant | Collision Resistant |
|---|---------------------------|----------------------------------|----------------------------|
| Hash + digital signature | yes | yes | yes* |
| Intrusion detection and virus detection | | yes | |
| Hash + symmetric encryption | | | |
| One-way password file | yes | | |
| MAC | yes | yes | yes* |

* Resistance required if attacker is able to mount a chosen message attack

Table 11.2 Hash Function Resistance Properties Required for Various Data Integrity Applications



Attacks on Hash Functions

► Brute-Force Attacks

- Does not depend on the specific algorithm, only depends on bit length
- In the case of a hash function, attack depends only on the bit length of the hash value
- Method is to pick values at random and try each one until a collision occurs

► Cryptanalysis

- An attack based on weaknesses in a particular cryptographic algorithm
- Seek to exploit some property of the algorithm to perform some attack other than an exhaustive search

Preimage and Second Preimage Attacks

- ▶ Adversary wishes to find a value y such that $H(y)$ is equal to a given hash value h
- ▶ The brute-force method is to pick values of y at random and try each value until a collision occurs
 - For an **m-bit** hash value, the level of effort is proportional to 2^m
 - The adversary would have to try, on average, 2^{m-1} values of y to find one that generates a given hash value h

Appendix E [Equation (E.1)]

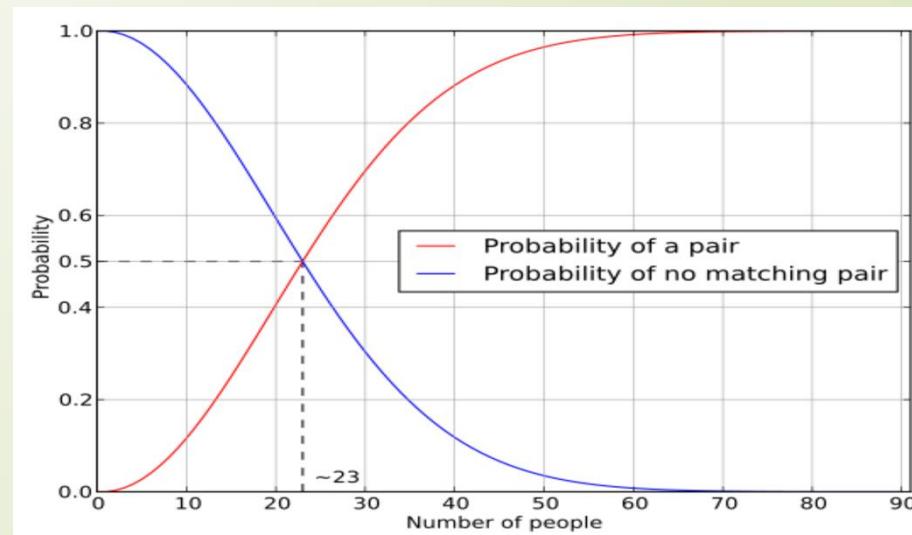


Collision Resistant Attacks

- ▶ Adversary wishes to find two messages or data blocks, x and y , that yield the same hash function: $H(x) = H(y)$
 - The effort required is explained by a mathematical result referred to as the birthday paradox
- ▶ Require considerably less effort than a preimage or second preimage attack
 - for an m -bit hash value, **if we pick data blocks at random**, we can expect to find two data blocks with the same hash value within $\sqrt{2^m} = 2^{m/2}$ Appendix E

Birthday Paradox

- The probability that, in a set of n randomly chosen people, at least two will share a birthday
 - Only 23 people are needed for that probability to exceed 50% (!)
 - If you are interested, read the calculations in [Wikipedia](#)
 - Also read this: [Birthday Attack](#)





Merkle – Damgard Structure

- ▶ Is the structure of most hash functions in use today, including SHA
- ▶ The hash function takes an input message and partitions it into **L fixed-sized blocks of b bits each**
 - If necessary, the final block is padded to b bits
 - The final block also includes the value of the total length of the input to the hash function
- ▶ The hash algorithm involves repeated use of a compression function, **f**
 - Takes two inputs (an **n-bit** input from the previous step, called the chaining variable and a **b-bit** block) and produces an **n-bit** output

Merkle – Damgard Structure

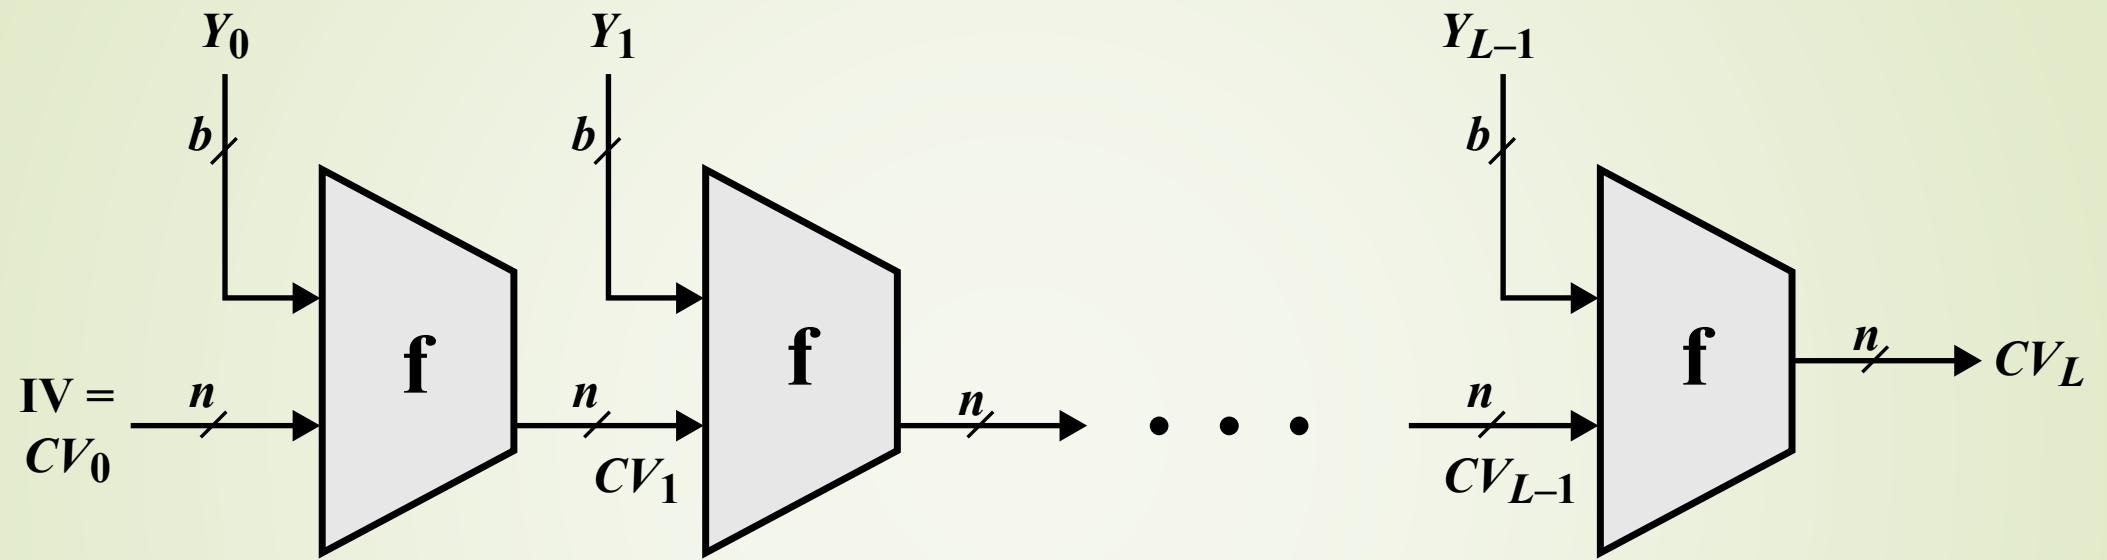
- The final value of the chaining variable is the hash value:

$$CV_0 = IV = \text{initial } n\text{-bit value}$$

$$CV_i = f(CV_{i-1}, Y_{i-1}) \quad 1 \leq i \leq L$$

$$H(M) = CV_L$$

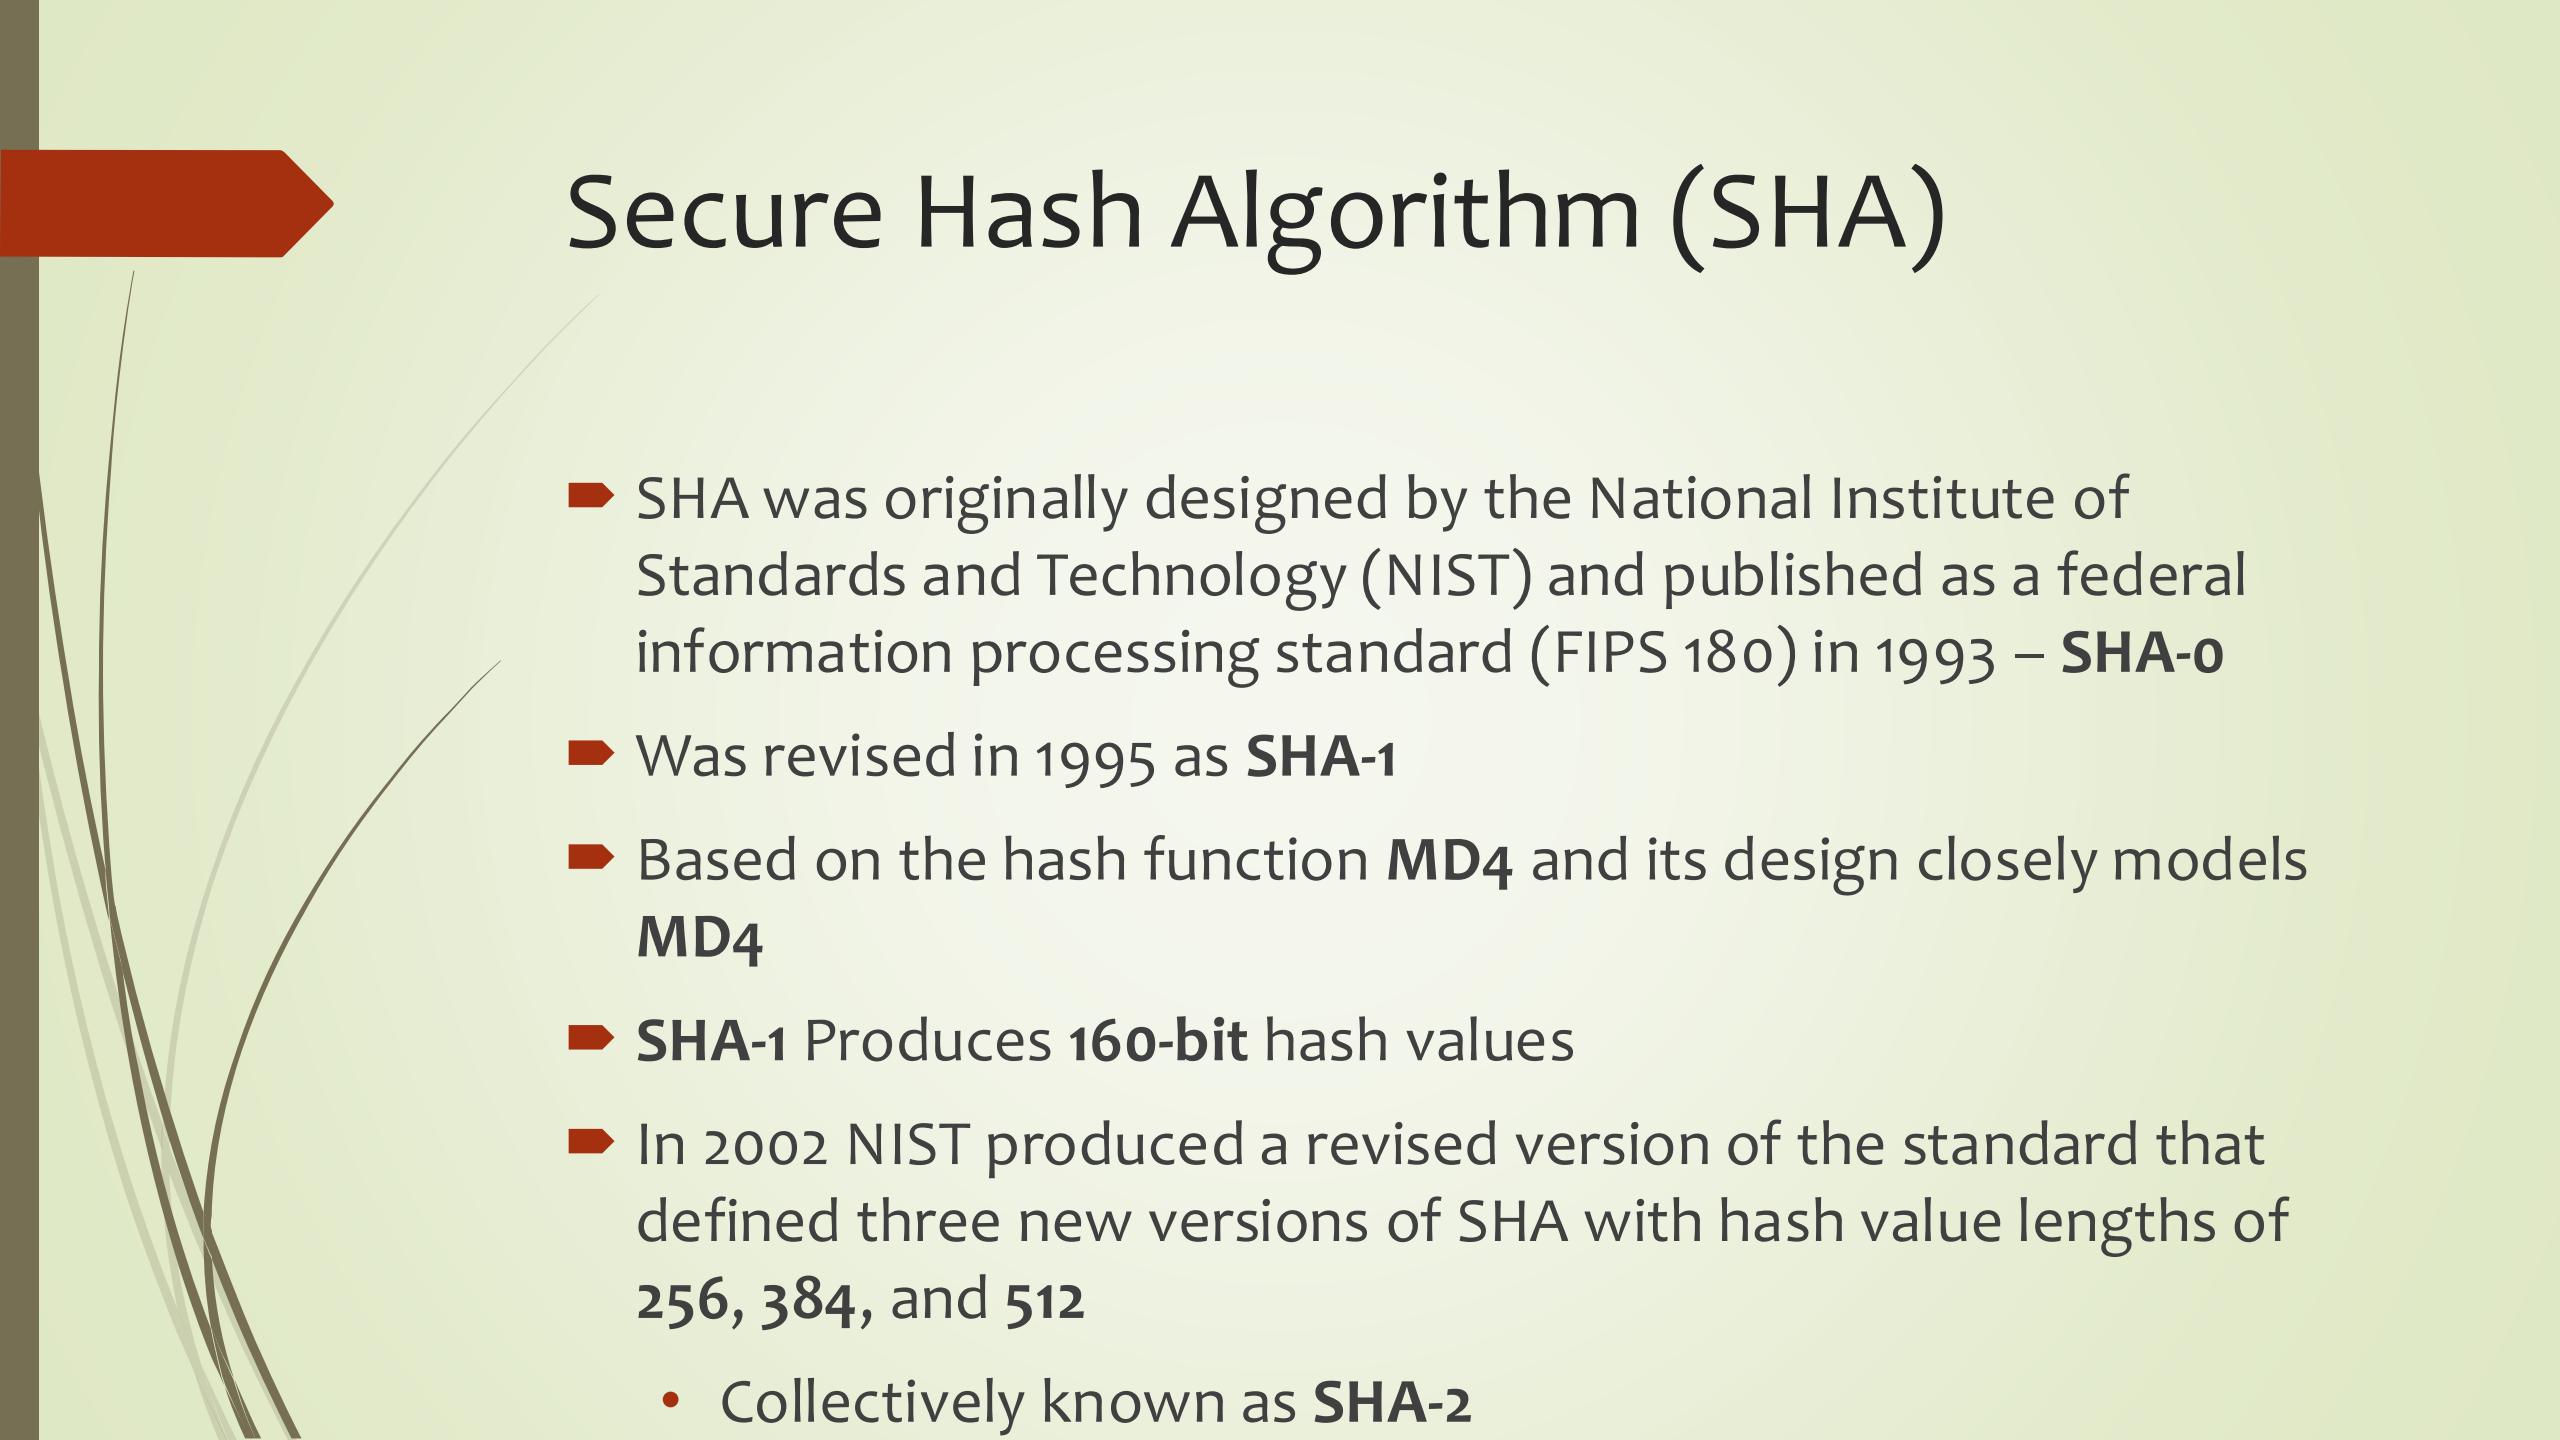
- The input to the hash function is a message M consisting of the blocks Y_0, Y_1, \dots, Y_{L-1}



IV = Initial value
 CV_i = chaining variable
 Y_i = i th input block
 f = compression algorithm

L = number of input blocks
 n = length of hash code
 b = length of input block

Figure 11.8 General Structure of Secure Hash Code – [Merkle-Damgård](#)



Secure Hash Algorithm (SHA)

- ▶ SHA was originally designed by the National Institute of Standards and Technology (NIST) and published as a federal information processing standard (FIPS 180) in 1993 – **SHA-0**
- ▶ Was revised in 1995 as **SHA-1**
- ▶ Based on the hash function **MD4** and its design closely models **MD4**
- ▶ **SHA-1** Produces **160-bit** hash values
- ▶ In 2002 NIST produced a revised version of the standard that defined three new versions of SHA with hash value lengths of **256, 384, and 512**
 - Collectively known as **SHA-2**



Secure Hash Algorithm (SHA)

- ▶ The most widely used hash function has been the Secure Hash Algorithm (SHA)
- ▶ A collision was found in **SHA-1** in 2017
- ▶ Microsoft, Google, Apple, and Mozilla have all announced that their respective browsers have stopped accepting SHA-1 SSL certificates in 2017

| Algorithm | Message Size | Block Size | Word Size | Message Digest Size |
|-------------|--------------|------------|-----------|---------------------|
| SHA-1 | $< 2^{64}$ | 512 | 32 | 160 |
| SHA-224 | $< 2^{64}$ | 512 | 32 | 224 |
| SHA-256 | $< 2^{64}$ | 512 | 32 | 256 |
| SHA-384 | $< 2^{128}$ | 1024 | 64 | 384 |
| SHA-512 | $< 2^{128}$ | 1024 | 64 | 512 |
| SHA-512/224 | $< 2^{128}$ | 1024 | 64 | 224 |
| SHA-512/256 | $< 2^{128}$ | 1024 | 64 | 256 |

* All sizes are measured in bits.

Table 11.3 Comparison of SHA Parameters

A SHA-256 Example

► Ashkan:

- 03cf25a5df5694becfd17663f774dd28e51d36c4ed4110373ac21a
64508dcc2a

► ashkan:

- e43adf78c5327baeb591e846de459826765963f82e41804f3cbf
eefcf706c0c1

► This Lecture PDF File:

- Try yourself!

Optional (FYI) Slides

For a Complete SHA-512 Implementation
refer to: [p. 353 to 361] of the Textbook

SHA-512

- The algorithm takes as input a **message with a maximum length** of less than 2^{128} bits and produces as **output a 512-bit message digest**
- 1. The input is processed in **1024-bit** blocks
 - The message is padded so that its length is congruent to 896 modulo 1024 [$\text{length} \equiv 896 \pmod{1024}$]
 - Padding is always added
 - Number of padding bits is in the range of 1 to 1024
 - Padding consists of a single 1 bit followed by the necessary number of 0 bits



SHA-512

2. A block of 128 bits is appended to the message
 - An unsigned 128-bit integer (most significant byte first)
 - contains the length of the original message in bits (before the padding)
3. A 512-bit buffer is used to hold intermediate and final results of the hash function
 - The buffer can be represented as eight 64-bit registers (a, b, c, d, e, f, g, h)

| | |
|----------------------|----------------------|
| a = 6A09E667F3BCC908 | e = 510E527FADE682D1 |
| b = BB67AE8584CAA73B | f = 9B05688C2B3E6C1F |
| c = 3C6EF372FE94F82B | g = 1F83D9ABFB41BD6B |
| d = A54FF53A5F1D36F1 | h = 5BE0CD19137E2179 |



SHA-512

4. The heart of the algorithm is a module that consists of **80 rounds**
 - Process message in 1024-bit (128-byte) blocks
 - Each round takes as input the 512-bit buffer value (abcdefg)
 - Each round t makes use of a 64-bit value W_t , derived from the current 1024-bit block being processed (M_i)

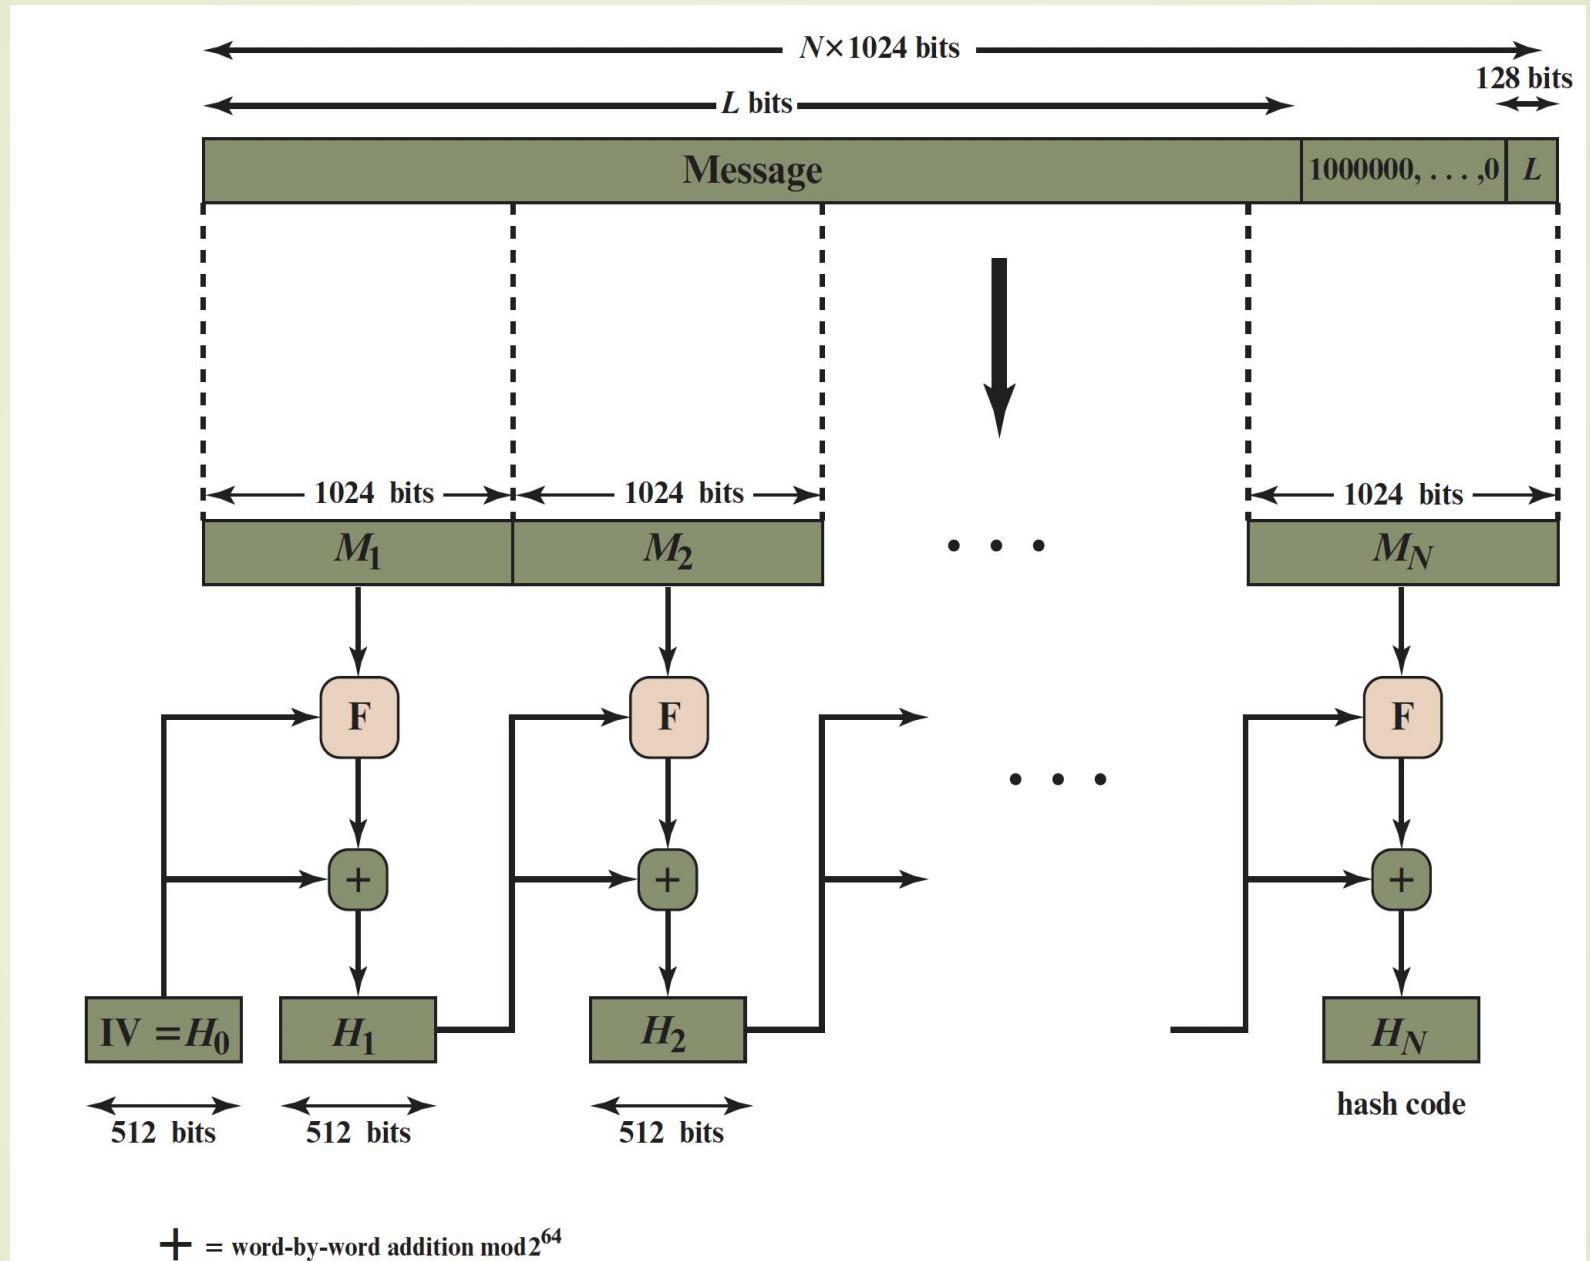


Figure 11.9 Message Digest Generation Using SHA-512

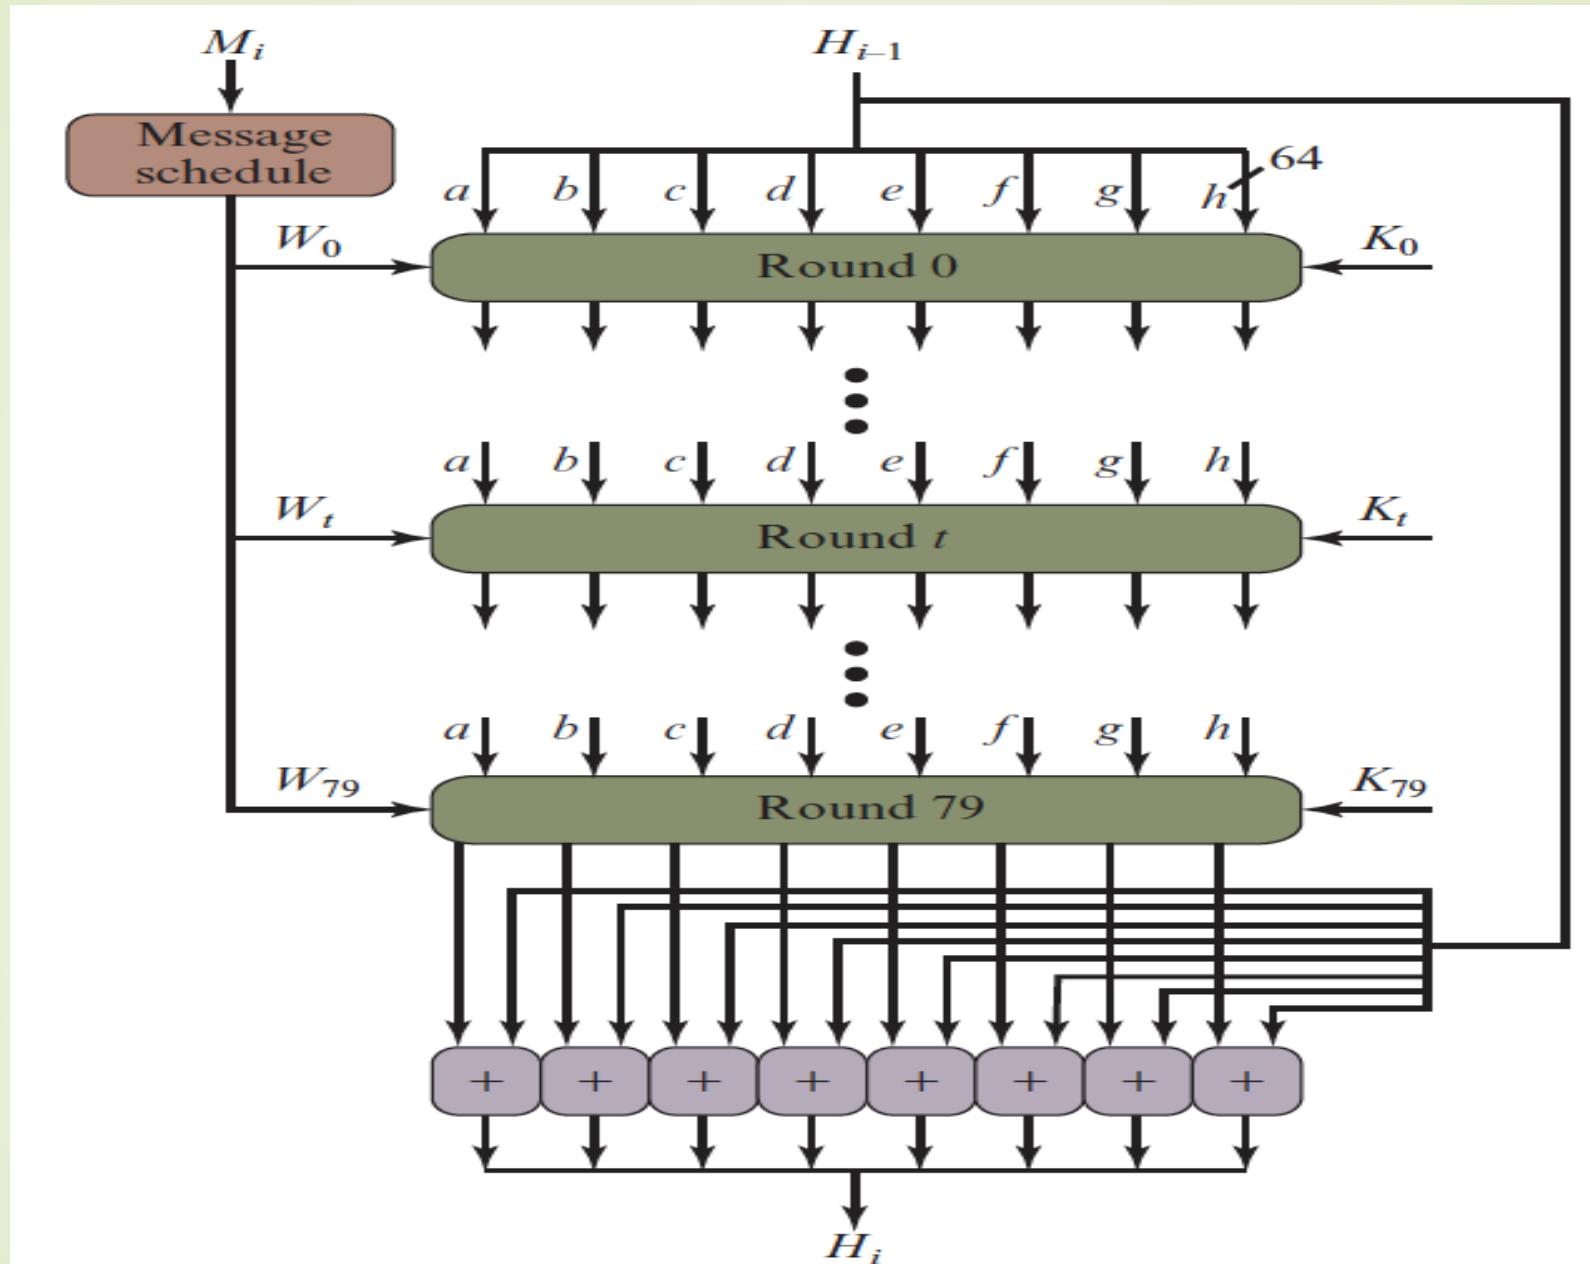


Figure 11.10 SHA-512 Processing of a Single 1024-Bit Block

| | | | |
|-------------------|------------------|-------------------|------------------|
| 428a2f98d728ae22 | 7137449123ef65cd | b5c0fbcfec4d3b2f | e9b5dba58189dbbc |
| 3956c25bf348b538 | 59f111f1b605d019 | 923f82a4af194f9b | ab1c5ed5da6d8118 |
| d807aa98a3030242 | 12835b0145706fbe | 243185be4ee4b28c | 550c7dc3d5ffb4e2 |
| 72be5d74f27b896f | 80deb1fe3b1696b1 | 9bdc06a725c71235 | c19bf174cf692694 |
| e49b69c19ef14ad2 | efbe4786384f25e3 | 0fc19dc68b8cd5b5 | 240ca1cc77ac9c65 |
| 2de92c6f592b0275 | 4a7484aa6ea6e483 | 5cb0a9dcbd41fdb4 | 76f988da831153b5 |
| 983e5152ee66dfab | a831c66d2db43210 | b00327c898fb213f | bf597fc7beef0ee4 |
| c6e00bf33da88fc2 | d5a79147930aa725 | 06ca6351e003826f | 142929670a0e6e70 |
| 27b70a8546d22fffc | 2e1b21385c26c926 | 4d2c6dfc5ac42aed | 53380d139d95b3df |
| 650a73548baf63de | 766a0abb3c77b2a8 | 81c2c92e47edaee6 | 92722c851482353b |
| a2bfe8a14cf10364 | a81a664bbc423001 | c24b8b70d0f89791 | c76c51a30654be30 |
| d192e819d6ef5218 | d69906245565a910 | f40e35855771202a | 106aa07032bbd1b8 |
| 19a4c116b8d2d0c8 | 1e376c085141ab53 | 2748774cdf8eeb99 | 34b0bcb5e19b48a8 |
| 391c0cb3c5c95a63 | 4ed8aa4ae3418acb | 5b9cca4f7763e373 | 682e6ff3d6b2b8a3 |
| 748f82ee5defb2fc | 78a5636f43172f60 | 84c87814a1f0ab72 | 8cc702081a6439ec |
| 90beffa23631e28 | a4506cebde82bde9 | bef9a3f7b2c67915 | c67178f2e372532b |
| ca273eceeaa26619c | d186b8c721c0c207 | eada7dd6cde0eb1e | f57d4f7fee6ed178 |
| 06f067aa72176fba | 0a637dc5a2c898a6 | 113f9804bef90dae | 1b710b35131c471b |
| 28db77f523047d84 | 32caab7b40c72493 | 3c9ebe0a15c9bebcb | 431d67c49c100d4c |
| 4cc5d4becb3e42b6 | 597f299cf657e2a | 5fc6fab3ad6faec | 6c44198c4a475817 |

Table 11.4 SHA-512 Constants

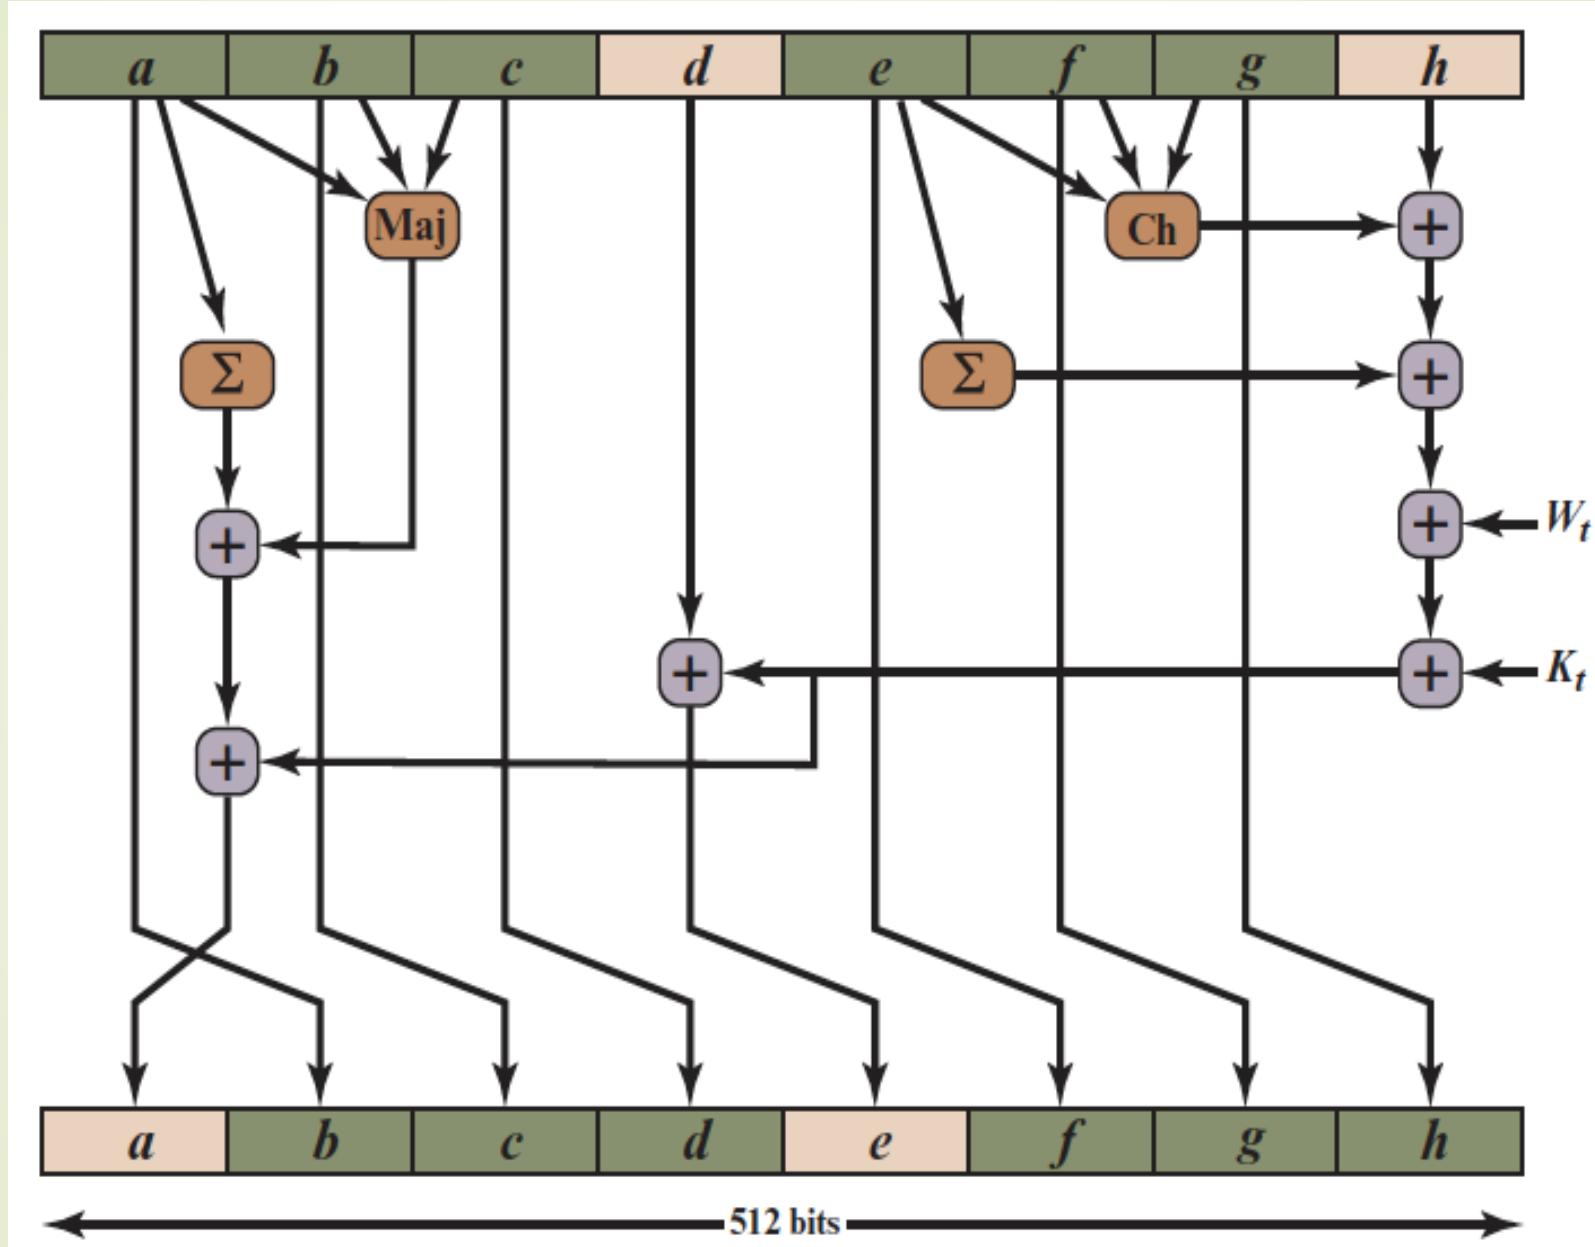


Figure 11.11 Elementary SHA-512 Operation (single round)



Real Life Example of SHA

- ▶ SHA-256 plays a fundamental role in blockchain technology
 - It generates the hash values of blocks in a blockchain, ensuring the immutability and integrity of the entire chain
- ▶ SHA-256 is the hash function and mining algorithm of the **Bitcoin** protocol
- ▶ Also used for transaction verification
- ▶ Bitcoin uses double SHA-256, meaning that it applies the hash functions twice
- ▶ <https://github.com/bitcoin/bitcoin>



Summary

- ▶ Applications of cryptographic hash functions
- ▶ Hash function used for Message Authentication needs to be secured
- ▶ Differences among preimage resistant, second preimage resistant, and collision resistant properties
- ▶ The basic structure of cryptographic hash functions

References

1. William Stallings: Cryptography and Network Security. Eight Edition 2020. ISBN-13: 9780135763971 Pearson
2. [Longitudinal Redundancy Check \(LRC\)/2-D Parity Check - Geeksforgeeks](#)
3. [The City University of New York](#)
4. [Birthday Paradox - University of Regina](#)
5. [The Birthday Paradox - Harvard CS125](#)
6. [The Birthday Paradox and Collisions in a Hash table – Medium - Arish Sateesan](#)
7. [Merkle-Damgård Revisited: how to Construct a Hash Function](#)